A gentle introduction to

Apache Spark, Databricks and Delta Lake

Personal Introduction

- Senior Consultant at Avanade
- MSc in Financial & Management Engineering
- Microsoft Certified Professional
- 12 years long journey, in the vast continent of Data

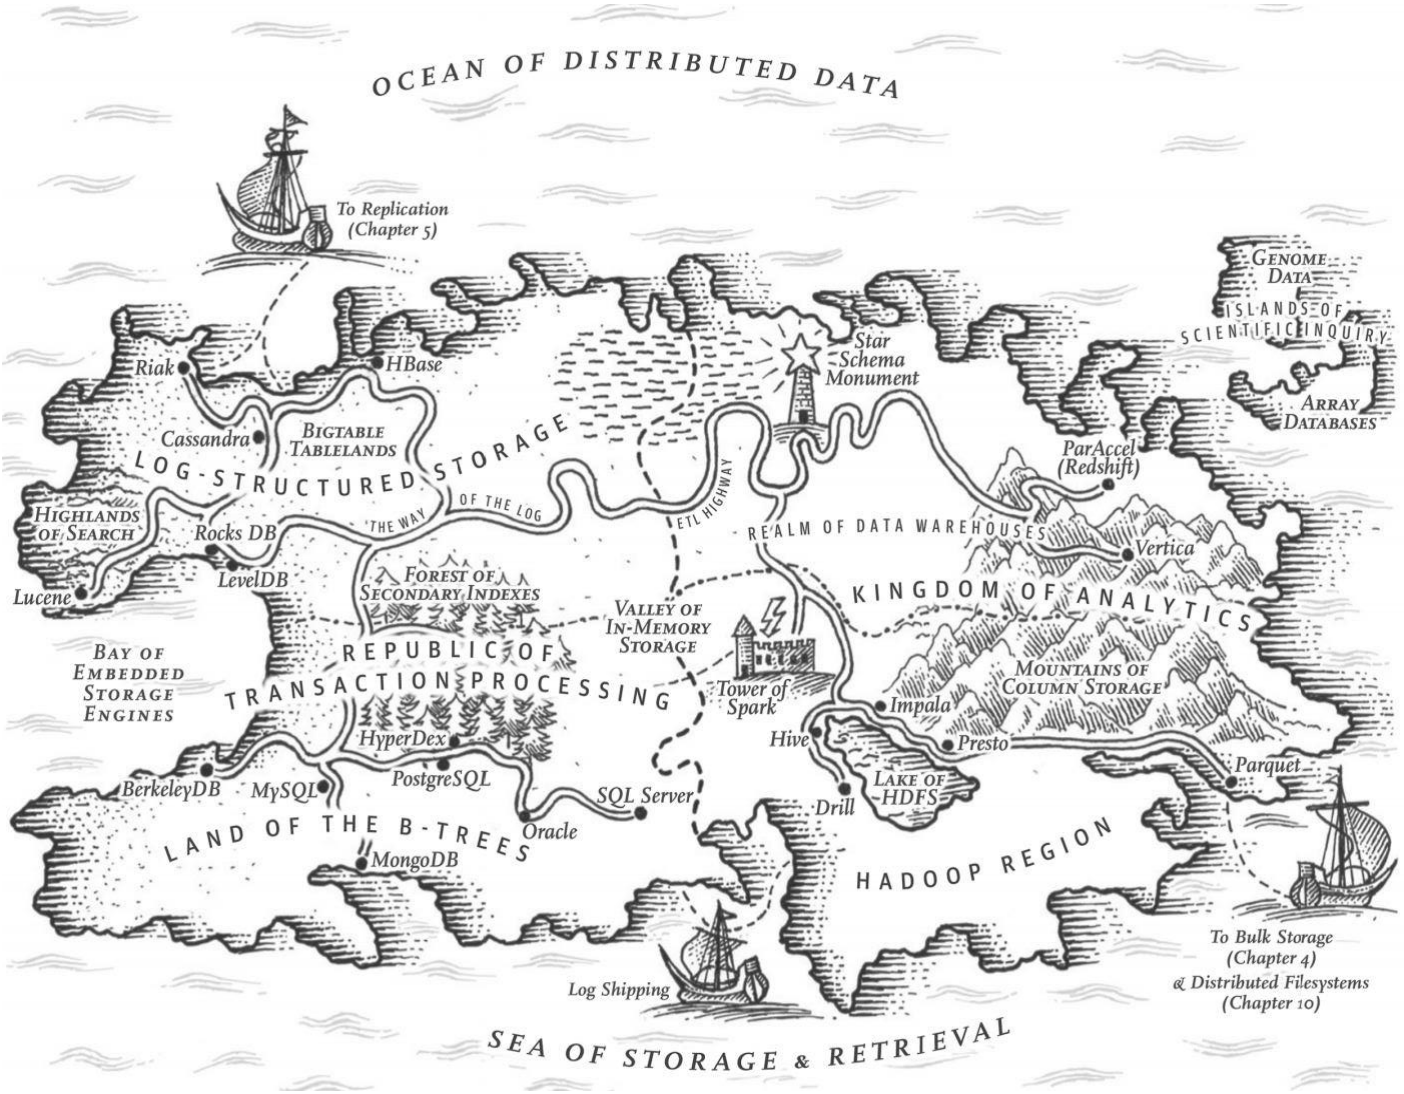

**Designing Data-Intensive Applications by Martin Kleppmann*

Agenda

- The big picture
- A peak into Apache Spark using Databricks
- Delta Lake
- Where do I go from here?

The evolution of Data Management

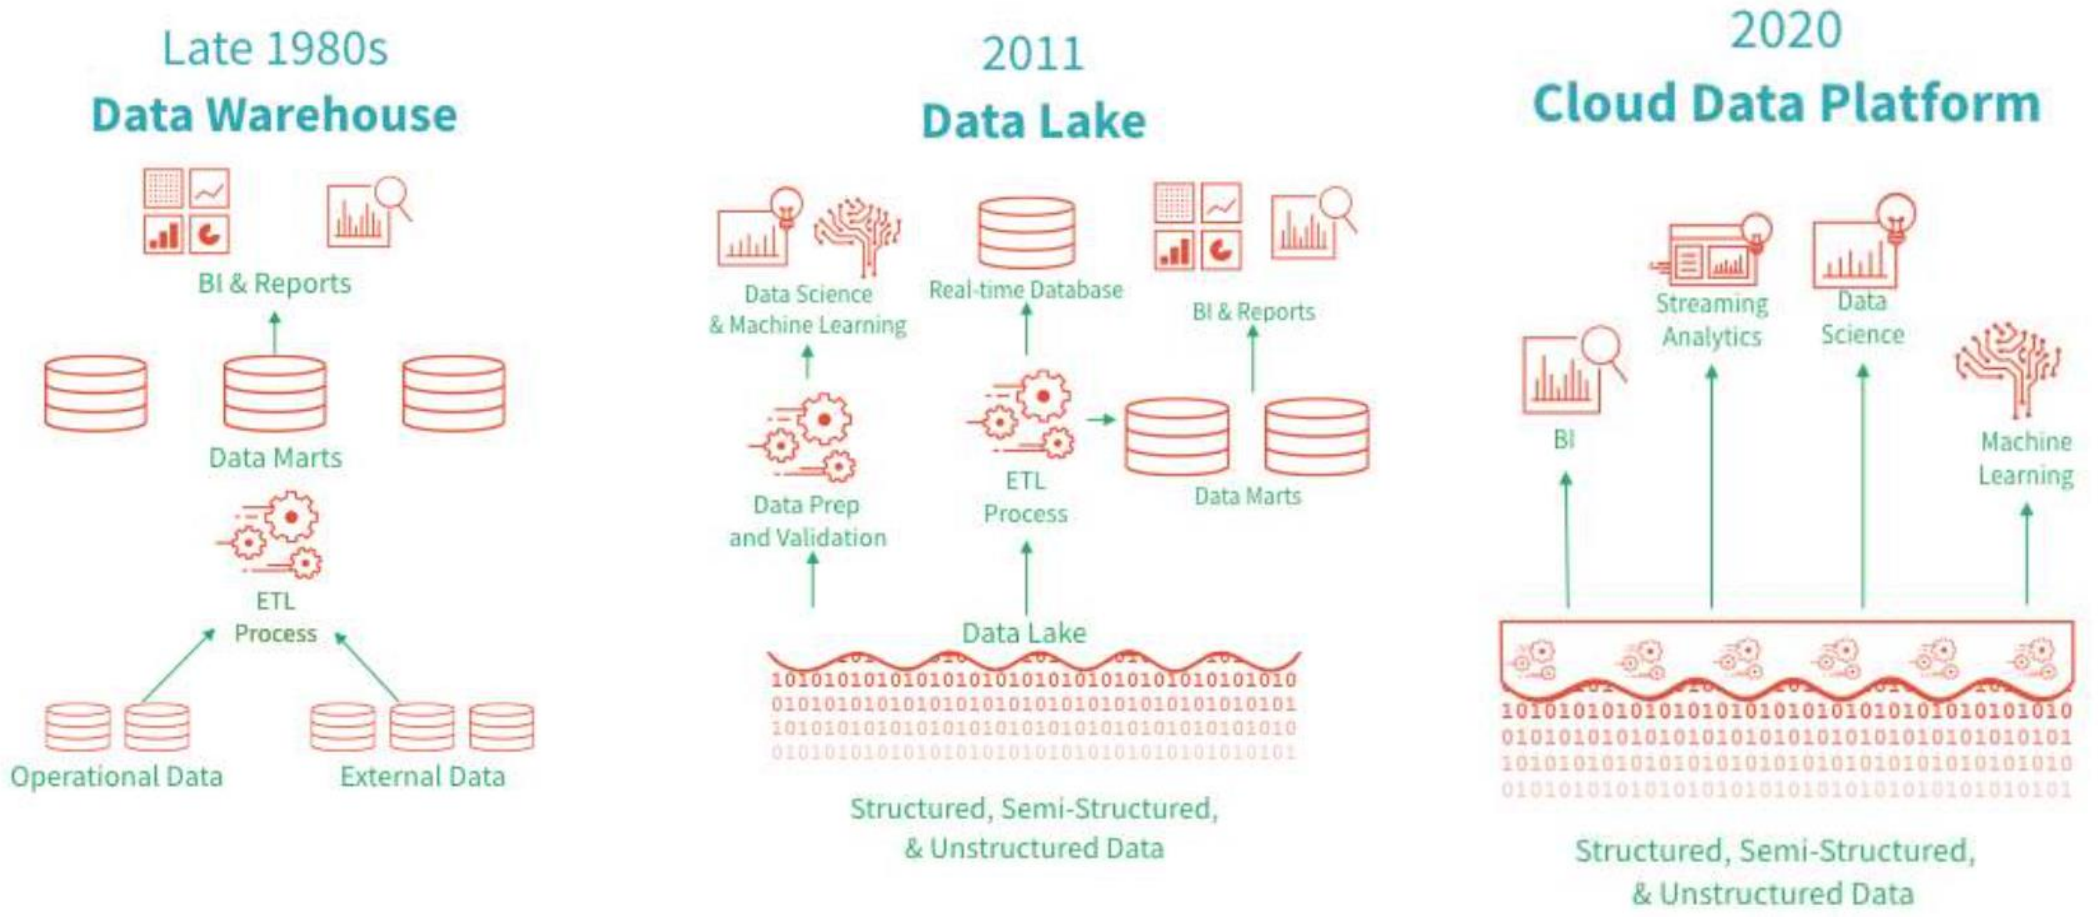

A brief history of Apache Spark

- Began in 2009 at UC Berkeley as a research project
- At that time Hadoop was the dominant paradigm
- Databricks was founded in 2013
- Spark 1.0 released in 2014
- Spark 2.0 released in 2016
- Spark 3.0 released in 2020

What is Spark and how does it work?

A simple example

Spreadsheet on a single machine

Table or Data Frame partitioned across servers in a data center

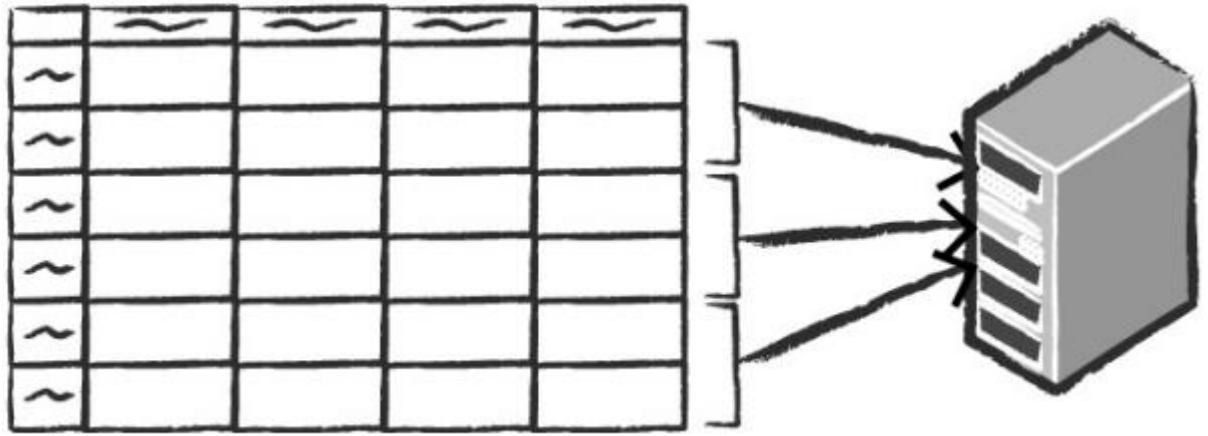

How do we talk to Spark?

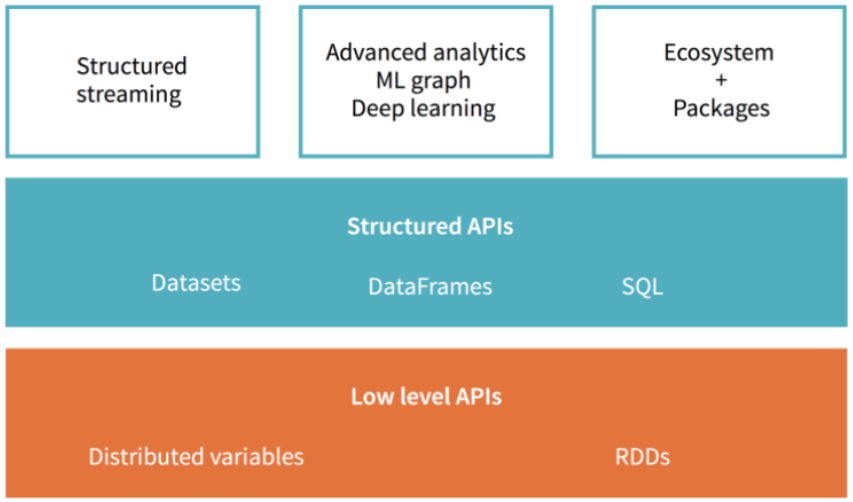

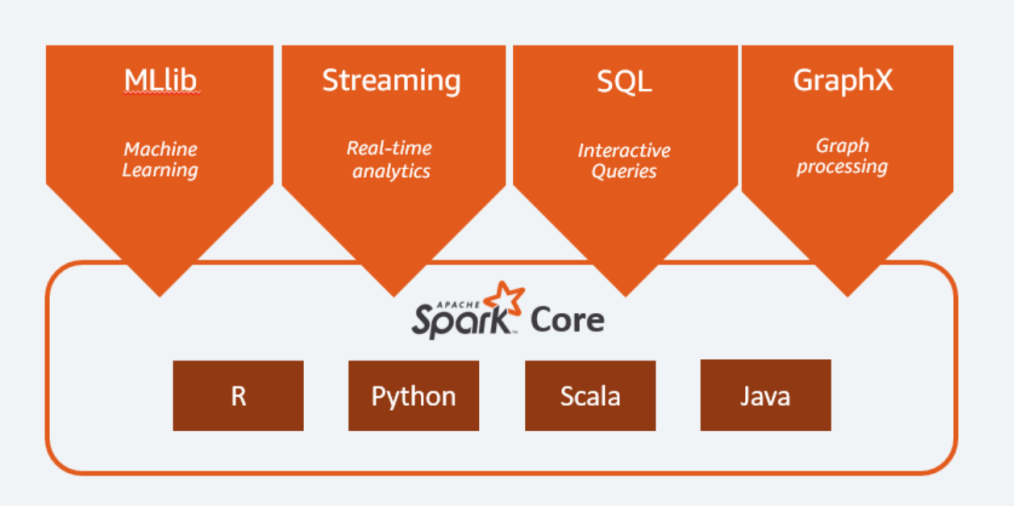

Distributing work

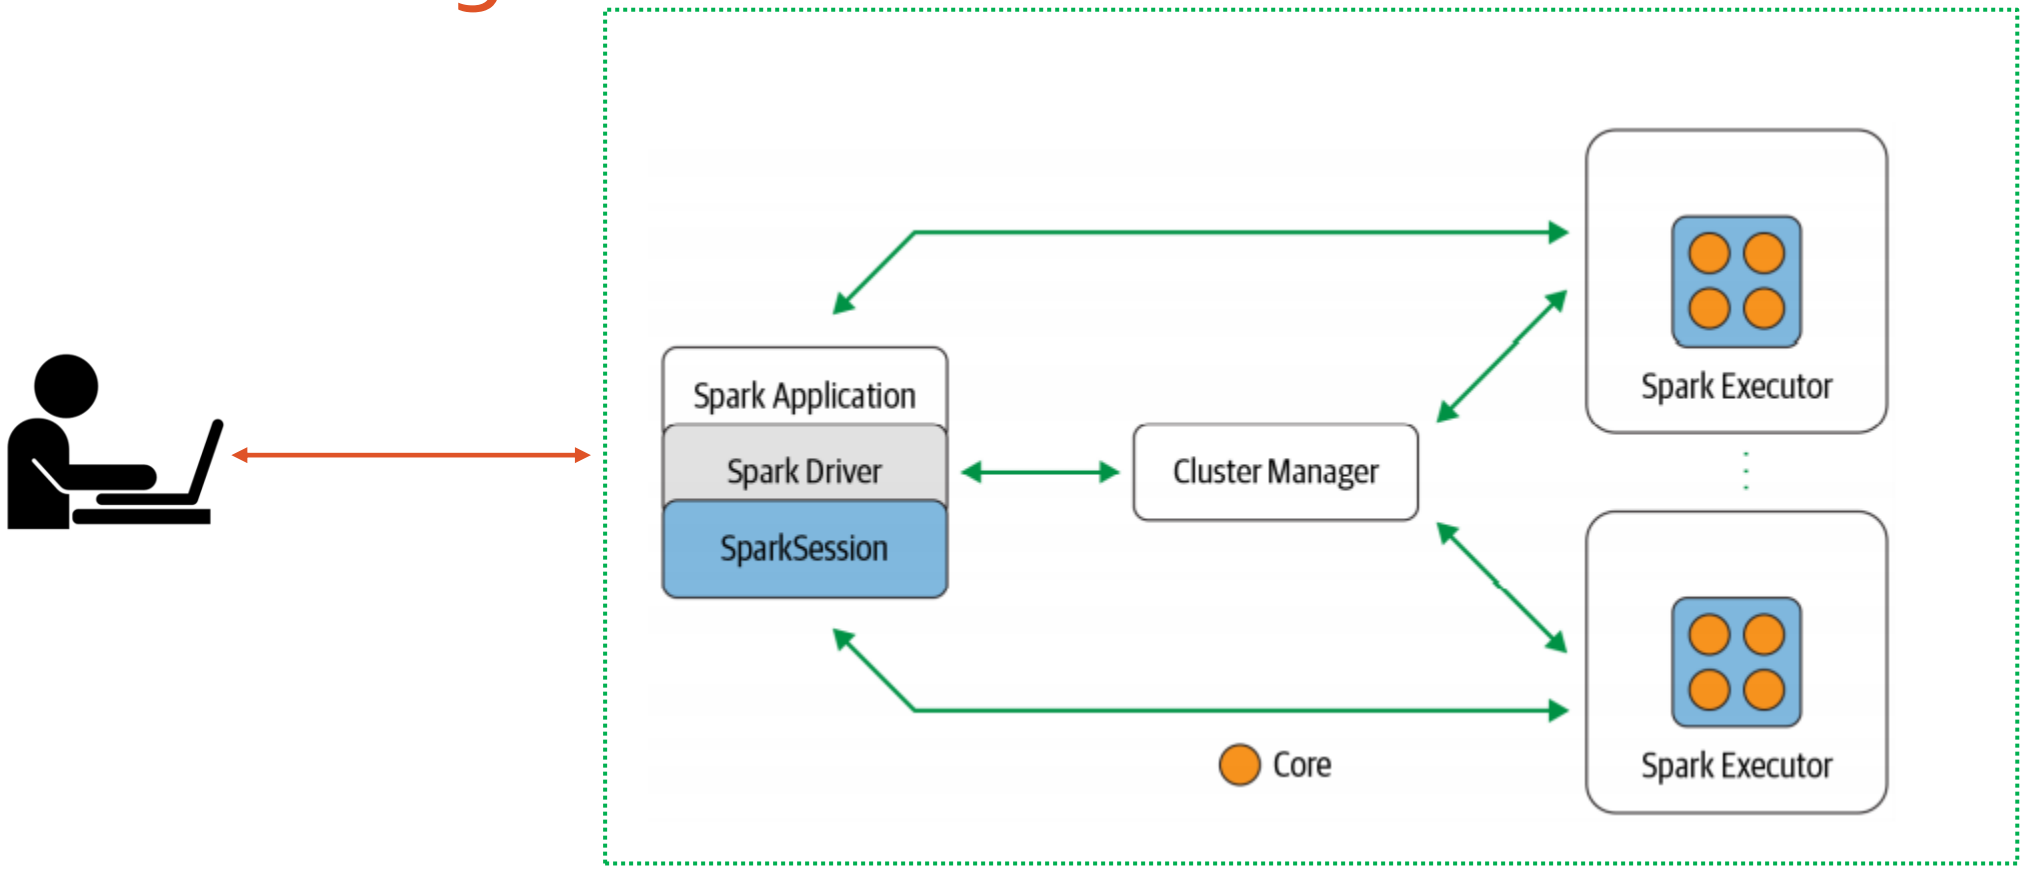

Distributing work

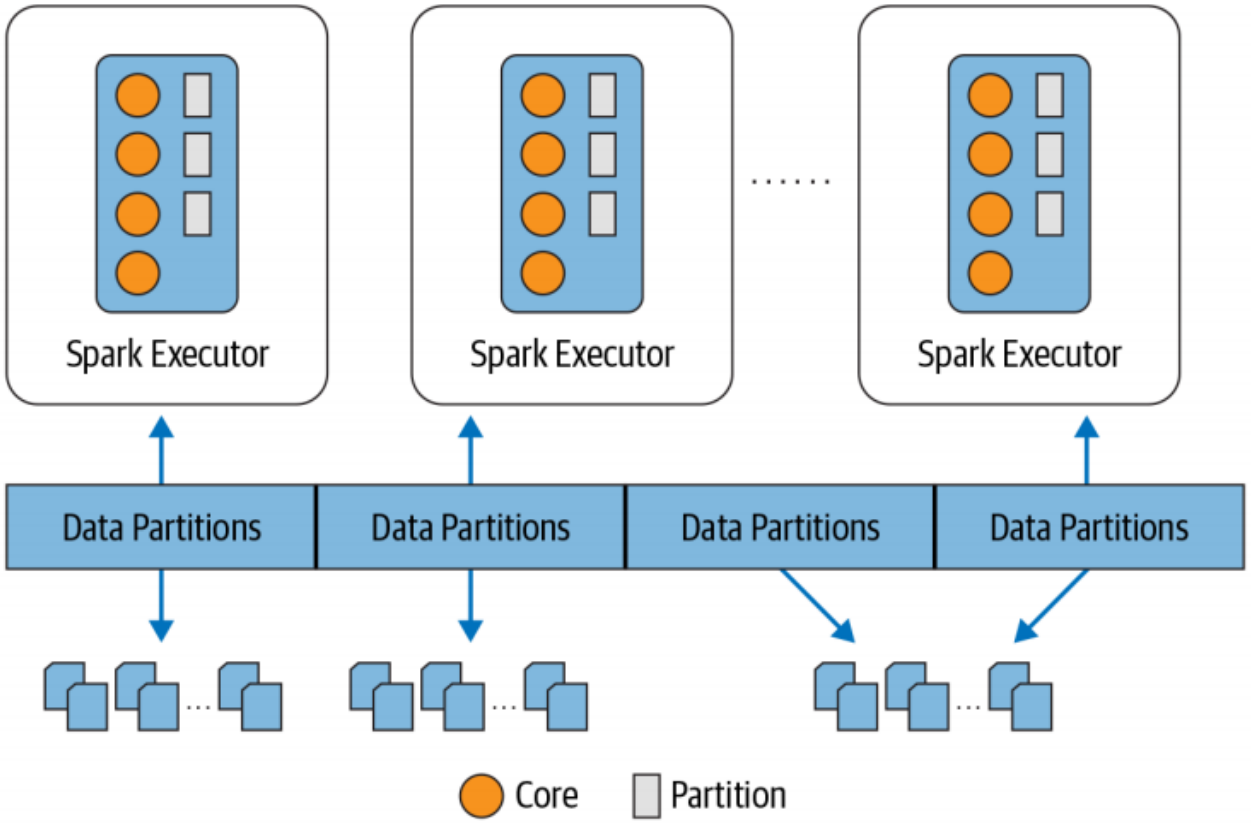

Spark stages, jobs and tasks

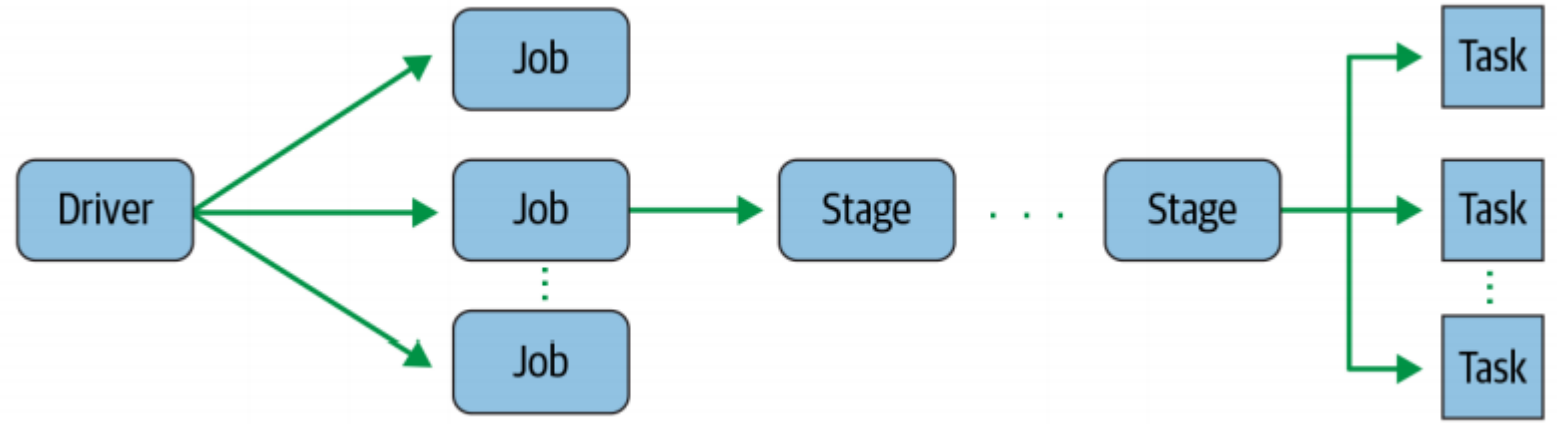

Transformations, Actions and Lazy Evaluation

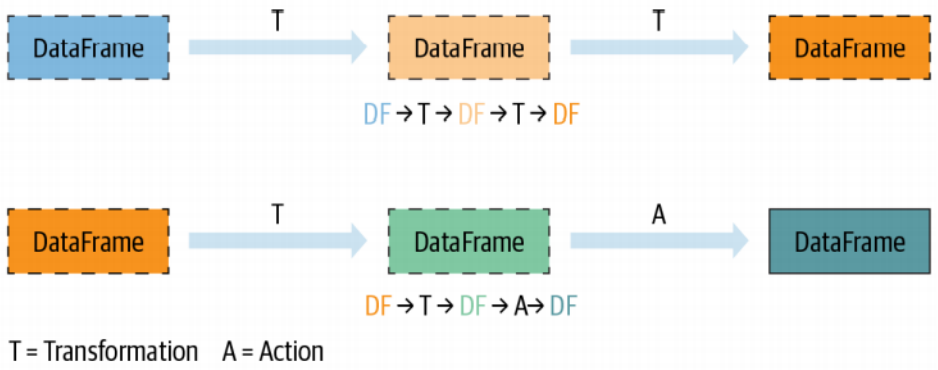


```python
# In Python
>>> strings = spark.read.text("../README.md")
>>> filtered = strings.filter(strings.value.contains("Spark"))
>>> filtered.count()
20
```

```scala
// In Scala
scala> import org.apache.spark.sql.functions._
scala> val strings = spark.read.text("../README.md")
scala> val filtered = strings.filter(col("value").contains("Spark"))
scala> filtered.count()
res5: Long = 20
```

Transformations, Actions and Lazy Evaluation

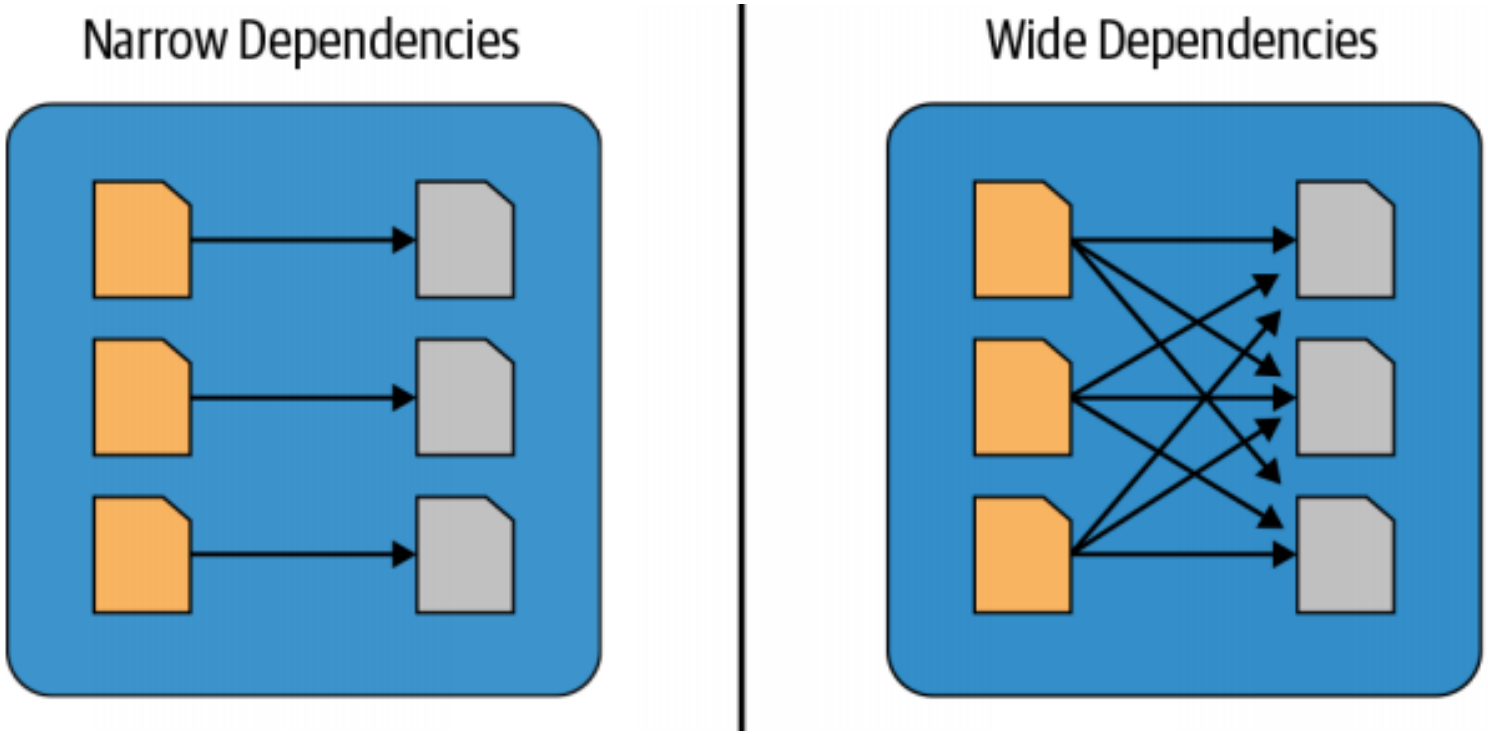

An example on narrow transformations:

"Select data where age = 37"

Node 1

Name	Age	City
Arnold	37	Amsterdam
Mohamed	37	London
John	25	Athens

Node 2

Name	Age	City
Lara	37	New York
George	31	London
Seif	45	Cairo

Node 3

Name	Age	City
Ankur	37	Mumbai
Jack	67	London
Lian	24	Beijing

Node 1

Name	Age	City
Arnold	37	Amsterdam
Mohamed	37	London

Node 2

Name	Age	City
Lara	37	New York

Node 3

Name	Age	City
Ankur	37	Mumbai

Final Result

Name	Age	City
Arnold	37	Amsterdam
Mohamed	37	London
Lara	37	New York
Ankur	37	Mumbai

Delta Lake

What is Delta Lake?

- Specifically designed to work with Apache Spark
- ACID compliant
- Streaming and batch unification
- Schema enforcement
- Time travel
- Upserts and deletes

Example Delta Lake architecture

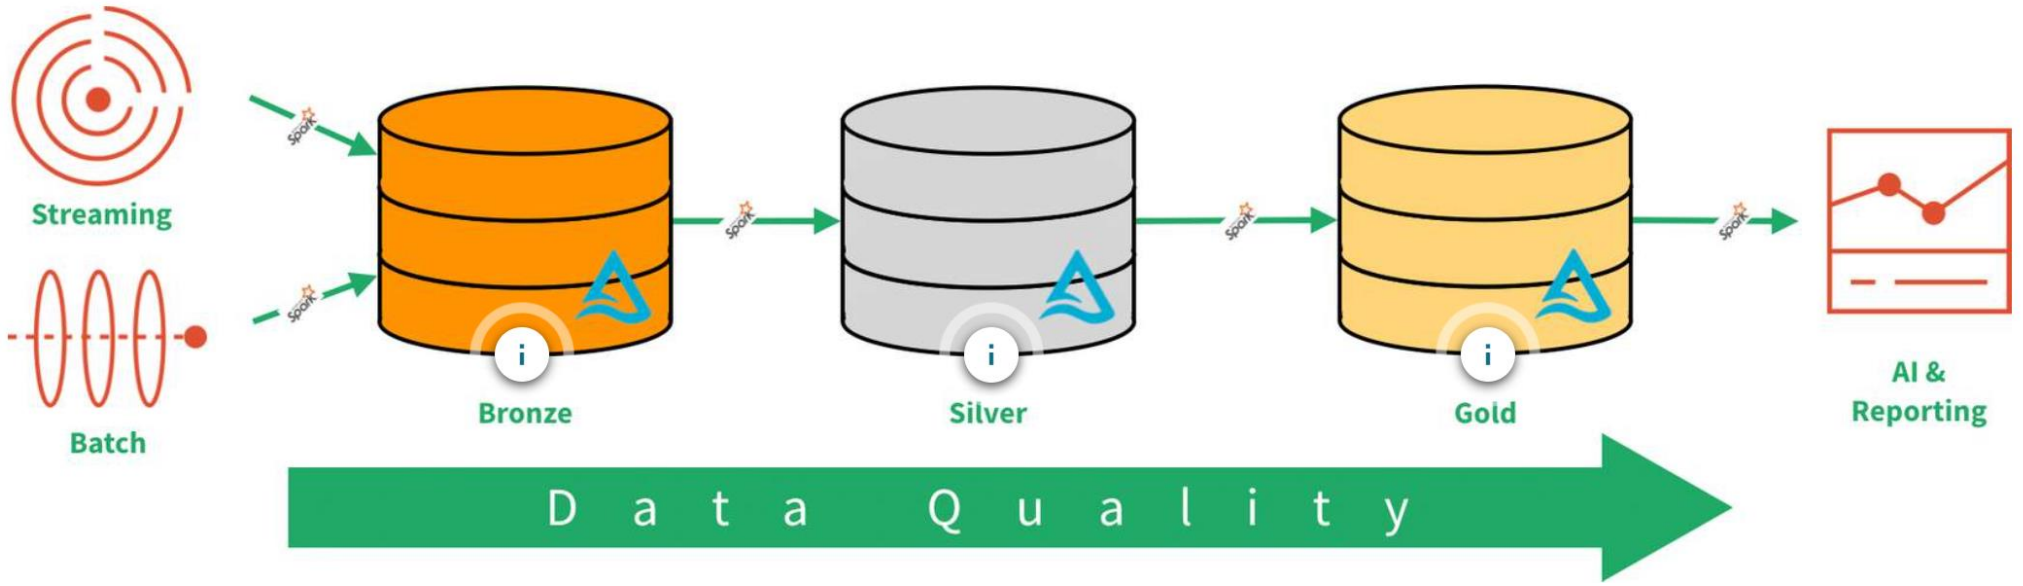

How can I use Spark?

- The ‘unmanaged way’
- The ‘managed way’

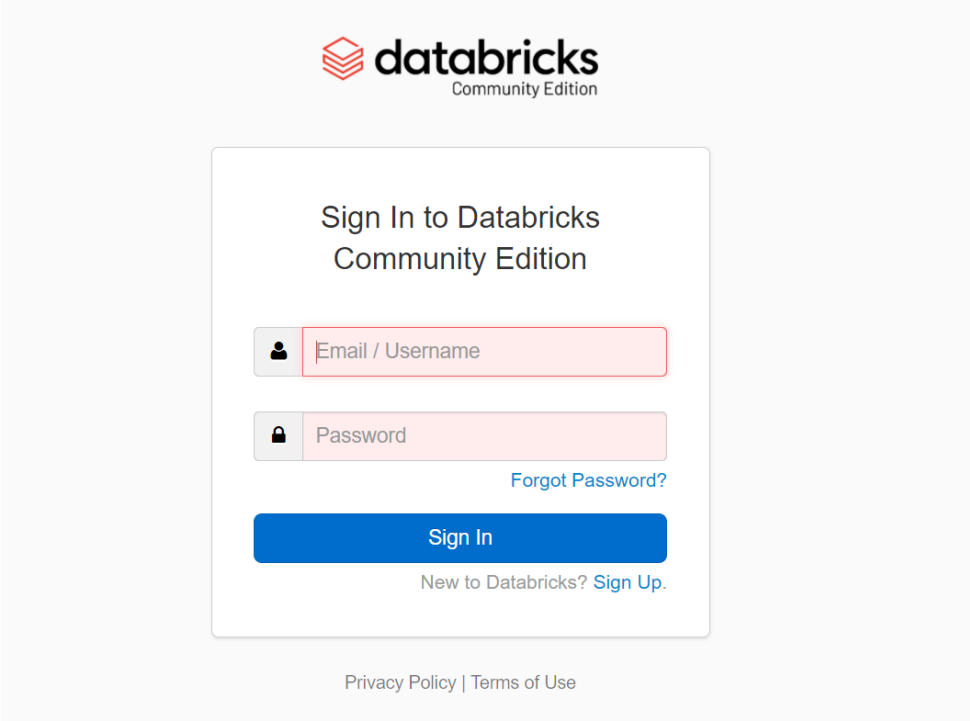

Demo

Source: Databricks Documentation

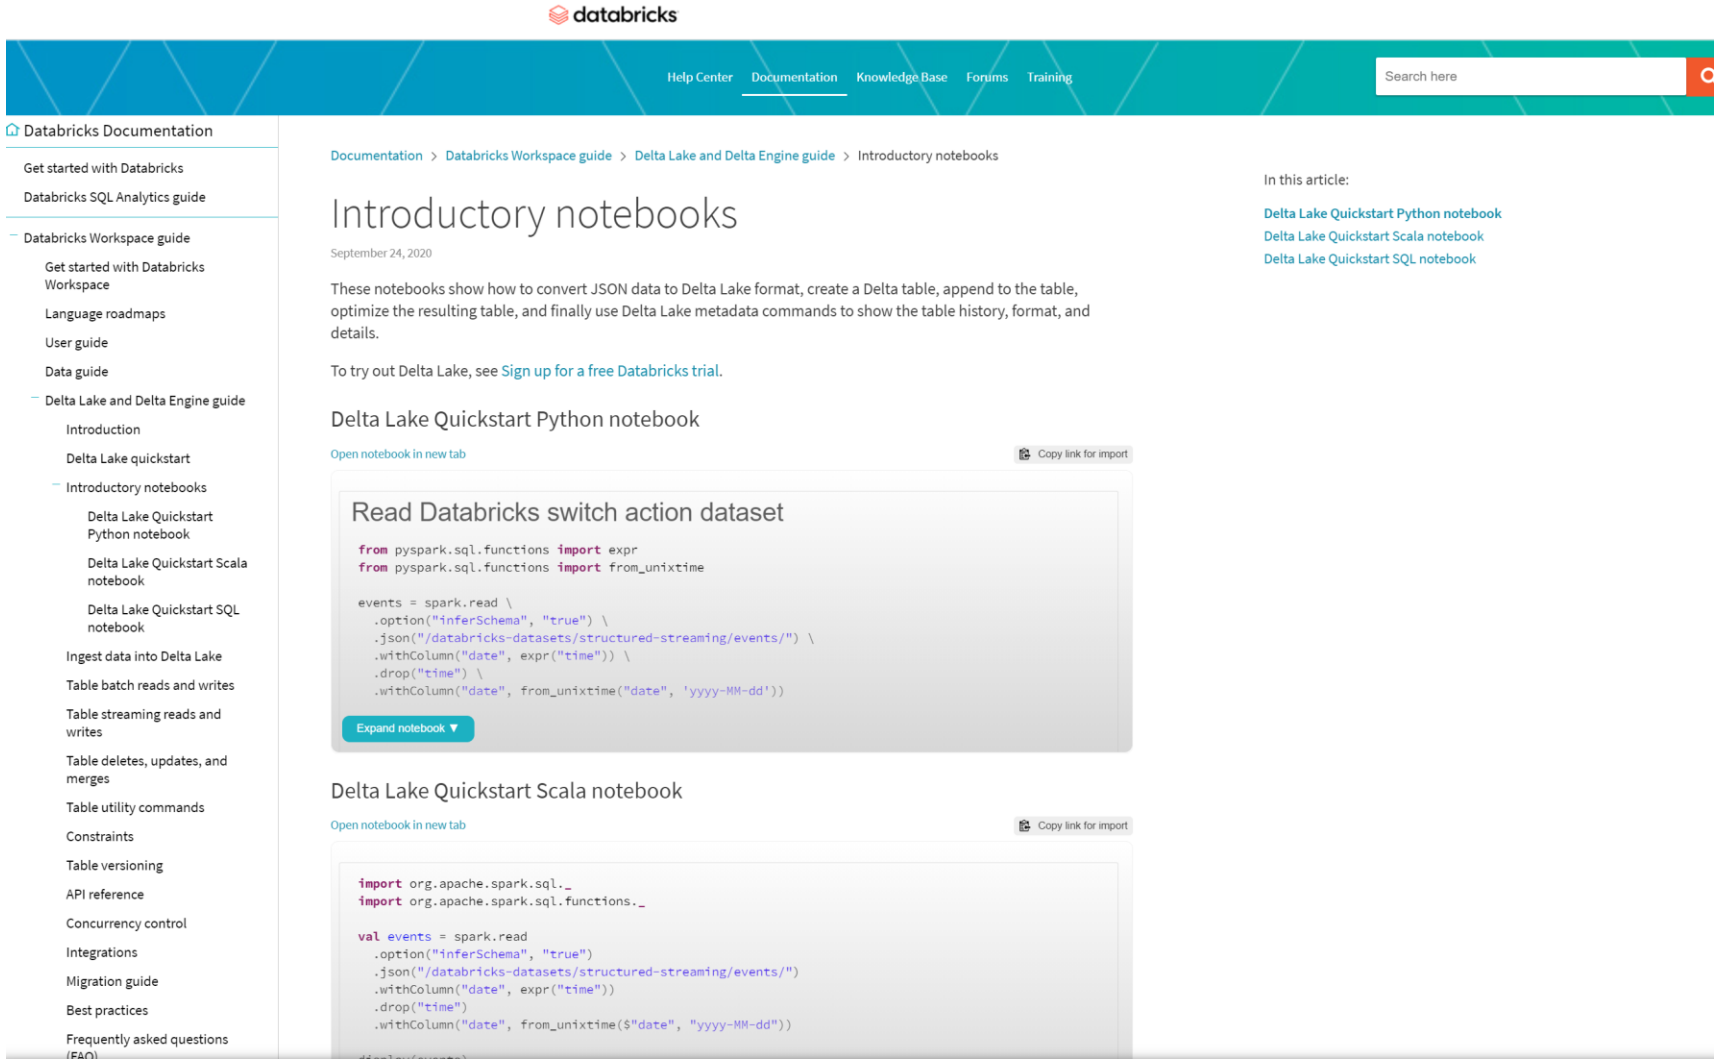

Where do I go from here

- https://www.oreilly.com/library/view/learning-spark-2nd/9781492050032/

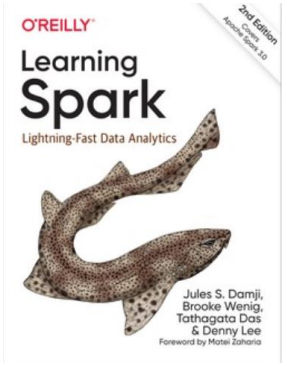

- https://academy.databricks.com/pathway/how-to-build-a-cloud-data-platform

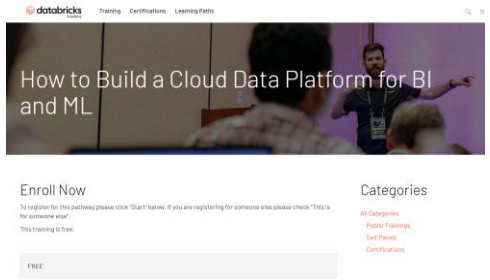

- https://docs.databricks.com/index.html

Thank you!

Any questions?